

# motor\_gpt\_dev

January 19, 2026

```
[1]: import torch
import torch.nn as nn
import torch.nn.functional as F
import random
import os, glob, math
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: device = "cuda" if torch.cuda.is_available() else "cpu"
print("device:", device)
```

device: cuda

## 0.0.1 Attention

```
[3]: class SelfAttentionHead(nn.Module):
    def __init__(self, head_dim, embed_dim, traj_size, dropout=0.0):
        super().__init__()
        # linear projections for Q, V, K
        self.key = nn.Linear(embed_dim, head_dim, bias=False)
        self.query = nn.Linear(embed_dim, head_dim, bias=False)
        self.value = nn.Linear(embed_dim, head_dim, bias=False)
        mask = torch.tril(torch.ones(traj_size, traj_size)).view(1, traj_size, traj_size)
        self.atten_drop = nn.Dropout(dropout)
        self.resid_drop = nn.Dropout(dropout)
        self.register_buffer("mask", mask)

    def forward(self, x):
        B, T, C = x.shape

        q = self.query(x) #(B, T, H)
        k = self.key(x)
        v = self.value(x)

        att = (q @ k.transpose(-2, -1)) * (1.0 / math.sqrt(k.size(-1))) # (B, T, T)
        att = att.masked_fill(self.mask[:, :T, :T]==0, float('-inf'))
```

```

        att = F.softmax(att, dim=-1) # (B, T, T)
        att = self.atten_drop(att)

        out = att @ v # (B, T, H)
        out = self.resid_drop(out)

    return out

```

[4]: # test the SelfAttentionHead

```

head = SelfAttentionHead(head_dim=16, embed_dim=32, traj_size=8)
x = torch.randn(4, 8, 32) # (B, T, C)
out = head(x)
print(out.shape)

```

`torch.Size([4, 8, 16])`

[5]: class MultiHeadAttention(nn.Module):

```

    def __init__(self, num_heads, embed_dim, head_dim, traj_size, dropout=0.0):
        super().__init__()
        self.heads = nn.ModuleList([SelfAttentionHead(head_dim, embed_dim, traj_size, dropout) for _ in range(num_heads)])
        self.proj = nn.Linear(num_heads * head_dim, embed_dim, bias=False)
        self.drop = nn.Dropout(dropout)

    def forward(self, x):
        multi_head_out = [h(x) for h in self.heads] # list of (B, T, head_size)
        multi_head_concat = torch.cat(multi_head_out, dim=-1) # (B, T, num_heads * head_size)

        out = self.drop(self.proj(multi_head_concat)) # (B, T, embed_dim)

    return out

```

[6]: mha = MultiHeadAttention(num\_heads=4, embed\_dim=32, head\_dim=8, traj\_size=8, dropout=0.0)

```

x = torch.randn(4, 8, 32)
out = mha(x)
print(out.shape)

```

`torch.Size([4, 8, 32])`

## 0.0.2 Transformer Block

```
[7]: class FeedForward(nn.Module):
    def __init__(self, embed_dim, expansion=4, dropout=0.0):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(embed_dim, expansion*embed_dim),
            nn.GELU(),
            nn.Linear(expansion*embed_dim, embed_dim),
            nn.Dropout(dropout),
        )
    def forward(self, x): return self.net(x)

class Block(nn.Module):
    def __init__(self, embed_dim, n_head, block_size, mlp_expansion=4, dropout=0.0):
        super().__init__()
        assert embed_dim % n_head == 0
        head_size = embed_dim // n_head
        self.ln1 = nn.LayerNorm(embed_dim)
        self.attn = MultiHeadAttention(n_head, embed_dim, head_size, block_size, dropout)
        self.ln2 = nn.LayerNorm(embed_dim)
        self.mlp = FeedForward(embed_dim, expansion=mlp_expansion, dropout=dropout)

    def forward(self, x):
        # TODO
        x = x + self.attn(self.ln1(x)) # skip connection
        x = x + self.mlp(self.ln2(x)) # skip connection

    return x
```

## 0.0.3 Motor GPT

```
[8]: class MotorGPT(nn.Module):
    def __init__(self, action_size=6, embed_dim=192, traj_size=128, n_layer=4, n_head=4, dropout=0.0):
        super().__init__()
        self.action_size = action_size
        self.traj_size = traj_size

        # "token embedding" for continuous actions
        self.token_emb = nn.Linear(action_size, embed_dim, bias=False)
        self.pos_emb = nn.Embedding(traj_size, embed_dim)

        self.blocks = nn.ModuleList([
```

```

        Block(embed_dim, n_head, traj_size, dropout=dropout)
        for _ in range(n_layer)
    ])
self.ln_f = nn.LayerNorm(embed_dim)

# regression head back to action space
self.head = nn.Linear(embed_dim, action_size, bias=True)

self.apply(self._init_weights)

def _init_weights(self, m):
    # NanoGPT init is fine conceptually, but add LayerNorm explicitly
    if isinstance(m, nn.Linear):
        nn.init.normal_(m.weight, mean=0.0, std=0.02)
        if m.bias is not None:
            nn.init.zeros_(m.bias)
    elif isinstance(m, nn.Embedding):
        nn.init.normal_(m.weight, mean=0.0, std=0.02)
    elif isinstance(m, nn.LayerNorm):
        nn.init.ones_(m.weight)
        nn.init.zeros_(m.bias)

def forward(self, x, targets=None):
    """
    x: (B, T, A) normalized actions
    targets: (B, T, A) next-step normalized actions
    """
    B, T, A = x.shape
    assert T <= self.traj_size, "Sequence length exceeds traj_size"

    #token and position embedding
    tok = self.token_emb(x) # (B, T, E)
    pos = self.pos_emb(torch.arange(T, device=x.device)).unsqueeze(0) # ↵ (1, T, E)
    h = tok + pos

    #pass through Transformer blocks
    for block in self.blocks:
        h = block(h)

    # regression head
    h = self.ln_f(h)
    pred = self.head(h)      # (B, T, A)

    loss = None
    if targets is not None:

```

```

        loss = F.mse_loss(pred, targets) # regression v1

    return pred, loss

@torch.no_grad()
def generate(self, seed_actions, max_new_steps=100, noise_std=0.0):
    """
    seed_actions: (B, T0, A) normalized actions
    returns:      (B, T0+max_new_steps, A) normalized actions
    """
    self.eval()
    out = seed_actions
    for _ in range(max_new_steps):
        cond = out[:, -self.traj_size:, :] # crop context, so only
    ↵last traj_size steps
        pred, _ = self(cond) # (B, Tcond, A), this
    ↵includes the whole context
        next_a = pred[:, -1, :] # (B, A), take only the
    ↵last time step

    # optional stochasticity (since MSE tends to be "average")
    if noise_std > 0:
        next_a = next_a + noise_std * torch.randn_like(next_a)

    out = torch.cat([out, next_a.unsqueeze(1)], dim=1)
    return out

```

```

[ ]: def estimate_loss(model, train_eps, val_eps, train_wlens, val_wlens,
                      trajectory_len, batch_size, eval_iters, device):
    # samples random batches from train and val sets, computes mean loss
    model.eval()
    out = {}
    with torch.no_grad():
        for name, eps, wl in [("train", train_eps, train_wlens), ("val", val_eps, val_wlens)]:
            losses = []
            for _ in range(eval_iters):
                xb, yb = get_batch_train(eps, wl, trajectory_len, batch_size, device)
                _, loss = model(xb, yb)
                losses.append(loss.item())
            out[name] = float(np.mean(losses))
    model.train()
    return out # returns {"train": mean_loss, "val": mean_loss}

def save_checkpoint(path, model, optimizer, step, best_val=None, metadata=None):

```

```

# builds a directory and saves model and optimizer state dicts
state = {
    "model": model.state_dict(),
    "optimizer": optimizer.state_dict(),
    "step": step,
    "best_val": best_val,
    "metadata": metadata or {},
}
torch.save(state, path)

def train_model(model, train_eps, val_eps, train_wlens, val_wlens,
               trajectory_len=128, batch_size=64, max_iters=2000,
               eval_interval=100, lr=3e-4, weight_decay=0.1,
               grad_clip=1.0, device="cuda", checkpoint_dir=None,
               save_interval=None, save_best=True, metadata=None):

    model.to(device)
    opt = torch.optim.AdamW(model.parameters(), lr=lr, weight_decay=weight_decay, betas=(0.9, 0.95))

    if checkpoint_dir is not None:
        os.makedirs(checkpoint_dir, exist_ok=True)
        if save_interval is None:
            save_interval = eval_interval

    best_val = float("inf")
    training_losses, validation_losses, eval_steps = [], [], []

    for it in range(max_iters):
        # sample batch, compute loss
        xb, yb = get_batch_train(train_eps, train_wlens, trajectory_len, batch_size, device)
        _, loss = model(xb, yb)

        # backprop and update
        opt.zero_grad(set_to_none=True)
        loss.backward()
        if grad_clip is not None:
            torch.nn.utils.clip_grad_norm_(model.parameters(), grad_clip)
        opt.step()

        # periodic evaluation
        if it % eval_interval == 0 or it == max_iters - 1:
            est = estimate_loss(model, train_eps, val_eps, train_wlens, val_wlens,

```

```

        trajectory_len, batch_size, eval_iters=25, u
device=device)
    print(f"iter {it:5d} | train {est['train']:.6f} | val {est['val']:.6f}")
    training_losses.append(est["train"])
    validation_losses.append(est["val"])
    eval_steps.append(it)

    if checkpoint_dir is not None:
        if save_best and est["val"] < best_val:
            best_val = est["val"]
            best_path = os.path.join(checkpoint_dir, "best.pt")
            save_checkpoint(best_path, model, opt, it, u
best_val=best_val, metadata=metadata)

        if it % save_interval == 0 or it == max_iters - 1:
            ckpt_path = os.path.join(checkpoint_dir, f"ckpt_{it:06d}.pt")
            save_checkpoint(ckpt_path, model, opt, it, u
best_val=best_val, metadata=metadata)

plt.figure()
plt.plot(eval_steps, training_losses, label="train")
plt.plot(eval_steps, validation_losses, label="val")
plt.xlabel("Iteration")
plt.ylabel("MSE loss")
plt.title("MotorGPT Train/Val Loss")
plt.legend()
plt.grid(True)
plt.show()

```